

# Deep Learning Group Challenge

## Hurricane Damage Detector

Jose Araujo, Mehran Chowdhury, Jonas Voßemer, Vincent Zitzewitz

23.06.2023


# What is the challenge about?

- Use satellite imagery data to detect damaged buildings after a hurricane

Damaged or not?



Damaged or not?



# What is the challenge about?

- Use a data set from a **paper by Cao and Choe (2020)**
  - Data provides to us includes **1358 images** (128x128) of training data, on which we applied a **80-20 training-validation** data split
  - **12228 test images** that can be used to assess the performance
  - The distribution of damages and no damages is **roughly balanced** in both
- 
- Goal defined for us: > **90% accuracy on test data**

# Why is this important? (Cao and Cheo, 2020)

- Damage assessment is critical
  - to emergency managers for efficient response
  - resource allocation
- How to gauge the damage extent, quantify the number of damaged buildings?
  - Traditionally done by ground survey
  - But this process can be labor-intensive and time-consuming
- **Can Deep Learning achieve high accuracy and replace the traditional process?**

# Approaches

- CNN model(s) from scratch (Jonas and Jose)
  - We used the original paper by Chao and Choe (2020) as a starting point
- Transfer Learning
  - Use Resnet50 as a base model and apply transfer learning (Mehran)
  - Use VGG16 as a base model and apply transfer learning (Vincent)

# CNN models from scratch

This is the “best” architecture used by Cao and Choe (2020) → CC2020

**Table 1** Convolutional neural network architecture that achieves the best result.

Layer type	Output shape	Number of trainable parameters
Input	3@(150x150)	0
2-D Convolutional 32@(3x3)	32@(148x148)	896
2-D Max pooling (2x2)	32@(74x74)	0
2-D Convolutional 64@(3x3)	64@(72x72)	18,496
2-D Max pooling (2x2)	64@(36x36)	0
2-D Convolutional 128@(3x3)	128@(34x34)	73,856
2-D Max pooling (2x2)	128@(17x17)	0
2-D Convolutional 128@(3x3)	128@(15x15)	147,584
2-D Max pooling (2x2)	128@(7x7)	0
Flattening	1x6272	0
Dropout	1x6272	0
Fully connected layer	1x512	3,211,776
Fully connected layer	1x1	513

*Note:* The total number of trainable parameters is 3,453,121.  $C@(A \times B)$  is interpreted as that there are a total of  $C$  matrices of shape  $(A \times B)$  stacked on top of one another to form a three-dimensional tensor. 2-D Max pooling layer with  $(2 \times 2)$  pooling size means that the input tensor's size will be reduced by a factor of 4.

In addition they applied

- Data augmentation
  - Rotation, horizontal flipping, ...
- Dropout before the fully connected layer was 50%
- L2 regularization in the fully connected layer with  $\lambda 10^{-6}$

# CNN models from scratch

## 1. CC2020 **without** data augmentation, dropout, and l2 regularization

Model: "sequential"

Layer (type)	Output Shape	Param #
rescaling (Rescaling)	(None, 128, 128, 3)	0
conv2d (Conv2D)	(None, 128, 128, 32)	896
max_pooling2d (MaxPooling2D)	(None, 64, 64, 32)	0
conv2d_1 (Conv2D)	(None, 64, 64, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 32, 32, 64)	0
conv2d_2 (Conv2D)	(None, 32, 32, 128)	73856
max_pooling2d_2 (MaxPooling2D)	(None, 16, 16, 128)	0
conv2d_3 (Conv2D)	(None, 16, 16, 128)	147584
max_pooling2d_3 (MaxPooling2D)	(None, 8, 8, 128)	0
flatten (Flatten)	(None, 8192)	0
dense (Dense)	(None, 512)	4194816
dense_1 (Dense)	(None, 1)	513

=====  
Total params: 4,436,161  
Trainable params: 4,436,161  
Non-trainable params: 0

	precision	recall	f1-score	support
no_damage	0.94	0.85	0.89	6333
damage	0.85	0.94	0.89	5895
accuracy			0.89	12228
macro avg	0.90	0.90	0.89	12228
weighted avg	0.90	0.89	0.89	12228

	0	1	
0	5390	943	
1	357	5538	

# CNN models from scratch

## 2. CC2020 **with** data augmentation, dropout and l2 regularization

- Trying to mimic CC2020 as close as possible in terms of architecture but also data augmentation and parameter choices

```
data_augmentation = tf.keras.Sequential([  
    layers.experimental.preprocessing.Rescaling(1./255, input_shape=(128, 128, 3)),  
    layers.experimental.preprocessing.RandomRotation(0.2),  
    layers.experimental.preprocessing.RandomFlip("horizontal"),  
    layers.experimental.preprocessing.RandomTranslation(0.1, 0.1),  
    layers.experimental.preprocessing.RandomZoom(0.2),  
])
```

	precision	recall	f1-score	support
no_damage	0.92	0.92	0.92	6333
damage	0.91	0.92	0.91	5895
accuracy			0.92	12228
macro avg	0.92	0.92	0.92	12228
weighted avg	0.92	0.92	0.92	12228

	0	1
0	5819	514
1	497	5398



# CNN models from scratch

## 3. CC2020 **with** data augmentation and dropout

- Reducing the number of convolution layers and increasing dense layers complexity

sequential (Sequential)	(None, 128, 128, 3)	0
conv2d (Conv2D)	(None, 126, 126, 32)	896
batch_normalization (Batch Normalization)	(None, 126, 126, 32)	128
max_pooling2d (MaxPooling2D)	(None, 63, 63, 32)	0
conv2d_1 (Conv2D)	(None, 61, 61, 64)	18496
batch_normalization_1 (Batch Normalization)	(None, 61, 61, 64)	256
max_pooling2d_1 (MaxPooling2D)	(None, 30, 30, 64)	0
conv2d_2 (Conv2D)	(None, 28, 28, 128)	73856
batch_normalization_2 (Batch Normalization)	(None, 28, 28, 128)	512
max_pooling2d_2 (MaxPooling2D)	(None, 14, 14, 128)	0
conv2d_3 (Conv2D)	(None, 12, 12, 128)	147584
batch_normalization_3 (Batch Normalization)	(None, 12, 12, 128)	512
max_pooling2d_3 (MaxPooling2D)	(None, 6, 6, 128)	0
flatten (Flatten)	(None, 4608)	0
dense (Dense)	(None, 512)	2359808
dropout (Dropout)	(None, 512)	0
dense_1 (Dense)	(None, 1)	513
Total params: 2,602,561		
Trainable params: 2,601,857		
Non-trainable params: 704		

Model: "sequential_2"		
Layer (type)	Output Shape	Param #
sequential (Sequential)	(None, 128, 128, 3)	0
conv2d_4 (Conv2D)	(None, 126, 126, 32)	896
max_pooling2d_4 (MaxPooling2D)	(None, 63, 63, 32)	0
conv2d_5 (Conv2D)	(None, 61, 61, 64)	18496
max_pooling2d_5 (MaxPooling2D)	(None, 30, 30, 64)	0
flatten_1 (Flatten)	(None, 57600)	0
dense_2 (Dense)	(None, 512)	29491712
dropout_1 (Dropout)	(None, 512)	0
dense_3 (Dense)	(None, 512)	262656
dropout_2 (Dropout)	(None, 512)	0
dense_4 (Dense)	(None, 256)	131328
dropout_3 (Dropout)	(None, 256)	0
dense_5 (Dense)	(None, 1)	257
Total params: 29,905,345		
Trainable params: 29,905,345		
Non-trainable params: 0		

# CNN models from scratch

## 3. CC2020 **with** data augmentation and dropout

- Reducing the number of convolution layers and increasing dense layers complexity

	precision	recall	f1-score	support
0	0.98	0.83	0.90	6333
1	0.84	0.98	0.90	5895
accuracy			0.90	12228
macro avg	0.91	0.90	0.90	12228
weighted avg	0.91	0.90	0.90	12228

---

Accuracy: 90.10%


---

Precision: 91.13%

---

Recall: 90.10%

---



	precision	recall	f1-score	support
0	0.80	0.88	0.84	6333
1	0.86	0.76	0.80	5895
accuracy			0.82	12228
macro avg	0.83	0.82	0.82	12228
weighted avg	0.83	0.82	0.82	12228

---

Accuracy: 82.26%

---

Precision: 82.62%

---


Recall: 82.26%

---

# CNN models from scratch

## 3. CC2020 **with** data augmentation and dropout

- Reducing the number of convolution layers and increasing dense layers complexity



# Transfer Learning: Resnet50

Resnet 50 model **with** data augmentation

```
conv5_block3_out (Activation) (None, 4, 4, 2048) 0 ['conv5_block3_add[0][0]']
global_average_pooling2d_16 (GlobalAveragePooling2D) (None, 2048) 0 ['conv5_block3_out[0][0]']
dense_47 (Dense) (None, 256) 524544 ['global_average_pooling2d_16[0][0]']
dense_48 (Dense) (None, 1) 257 ['dense_47[0][0]']

=====
Total params: 24,112,513
Trainable params: 524,801
Non-trainable params: 23,587,712
```

Confusion Matrix:

```
[[5978  355]
```

```
 [ 396 5499]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.94	0.94	0.94	6333
1	0.94	0.93	0.94	5895
accuracy			0.94	12228
macro avg	0.94	0.94	0.94	12228
weighted avg	0.94	0.94	0.94	12228

Accuracy: 0.9385835786719006

# Transfer Learning: Resnet50

Problems replicating results! Data augmentation?

	precision	recall	f1-score	support
0	0.78	0.82	0.80	6333
1	0.80	0.75	0.77	5895
accuracy			0.79	12228
macro avg	0.79	0.79	0.79	12228
weighted avg	0.79	0.79	0.79	12228

---

Accuracy: 78.75%


---

Precision: 78.80%

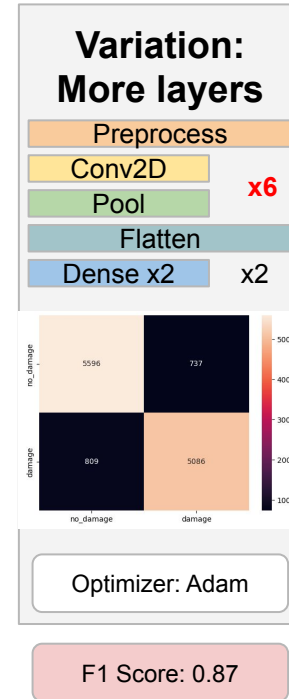
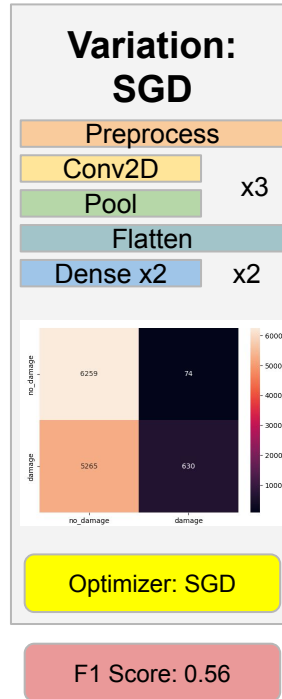
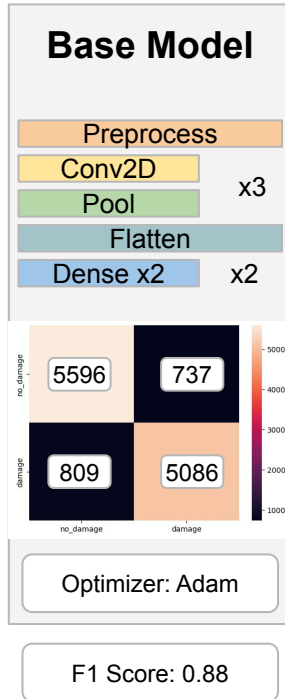
---

Recall: 78.75%



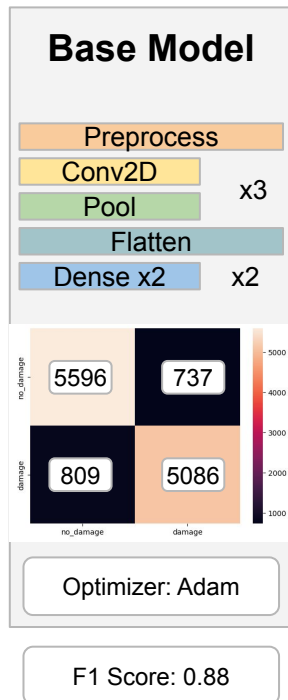
---





# Transfer Learning: VGG16



# Transfer Learning: VGG16



F1 Score: 0.91



F1 Score: 0.93

# Conclusion

- One needs to be extra careful when working with small dataset and large/complex models
- Data augmentation helped improving the performance, but also one needs to be careful
- Hard to replicate (original/colleagues) results using similar models